

# Multilayer Perceptron

---

Introduction à l'apprentissage automatique – GIF-4101 / GIF-7005

Professor: Christian Gagné

Week 7



UNIVERSITÉ  
LAVAL


## **7.1   Multilayer perceptron model**

---


# Natural intelligence

- Brain: natural intelligence
  - Parallel and distributed computing
  - Learning and generalization
  - Adaptation and context
  - Error-tolerant
  - Low energy consumption
- Biological computational machine!

# Biological neuron



## Artificial neuron model

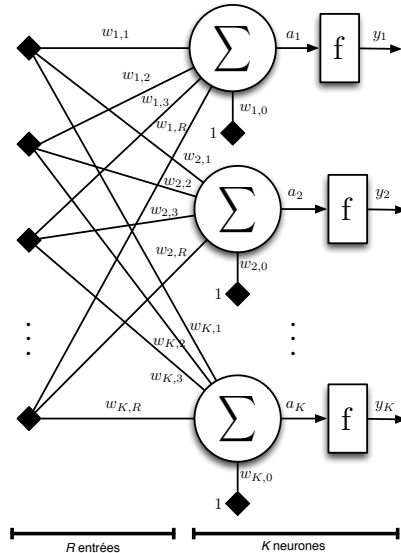


- Each neuron is a linear discriminator with a transfer function  $f$

$$y = f \left( \sum_i w_i x_i + w_0 \right) = f(\mathbf{w}^\top \mathbf{x} + w_0)$$

- Examples of transfer functions
  - Linear function:  $f_{lin}(a) = a$
  - Sigmoid function:  $f_{sig}(a) = \frac{1}{1+\exp(-a)}$
  - Step function:  $f_{step}(a) = 1$  if  $a \geq 0$  and  $f_{step}(a) = 0$  otherwise
- Several neurons connected together form a neural network
  - Single-layer network: neurons are connected to the inputs
  - Multilayer network: some neurons are connected to the outputs of other neurons

# Neural network (one layer)

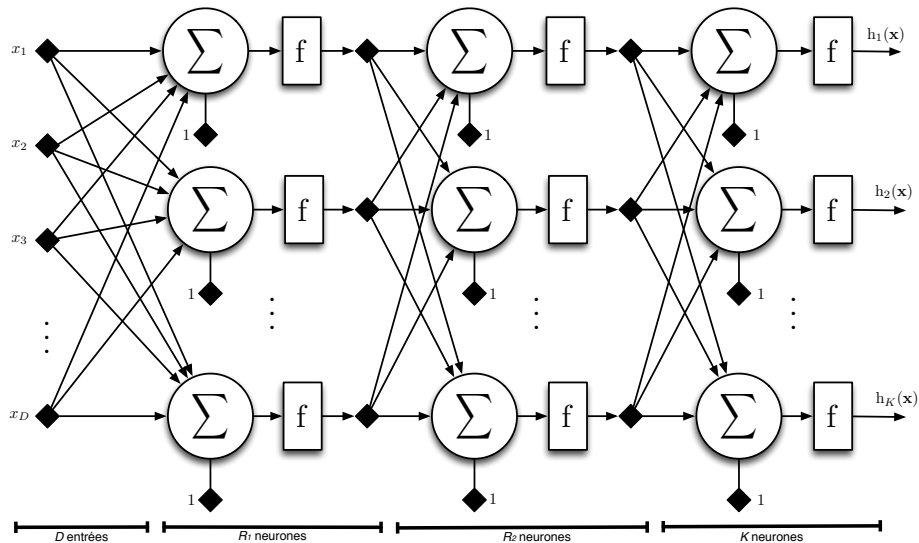


# Multilayer perceptron

- Single-layer network: set of linear discriminants
  - Unable to correctly classify non-linearly separable data
- Multilayer network (multilayer perceptron)
  - Linear discriminants (neurons) cascaded at the output of other linear discriminants
  - Able to classify non-linearly separable data
  - Set of simple classifiers
  - Each layer makes a projection into a new space
- During data processing, information is propagated from inputs to outputs



# Multilayer perceptron



## 7.2 Topology and capacity of networks

---

# XOR problem

- XOR problem


$$\mathbf{x}_1 = [0 \ 0]^\top \quad r_1 = 0$$

$$\mathbf{x}_2 = [0 \ 1]^\top \quad r_2 = 1$$



$$\mathbf{x}_3 = [1 \ 0]^\top \quad r_3 = 1$$

$$\mathbf{x}_4 = [1 \ 1]^\top \quad r_4 = 0$$

- Example of non-linearly separable data




# Network for the XOR problem




- Depending on the network topology used, different decision boundaries are possible
  - Network with a hidden layer and an output layer: convex boundaries
  - Two or more hidden layers: concave boundaries
    - The neural network is then a universal approximator
- Number of weights (therefore of neurons) directly determines the complexity of the classifier
  - Determining the right topology is often a matter of trial and error


# Types of decision boundaries




Convex open



Convex close

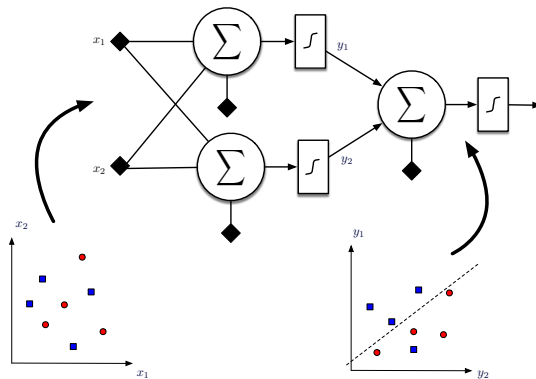


Concave open




Concave close

## Number of neurons on the hidden layer (classification)



2 neurons on the hidden layer: non-optimal

## Number of neurons on the hidden layer (classification)




3 neurons on the hidden layer: no error




## Number of neurons on the hidden layer (regression)


1 neuron



3 neurons



9 neurons



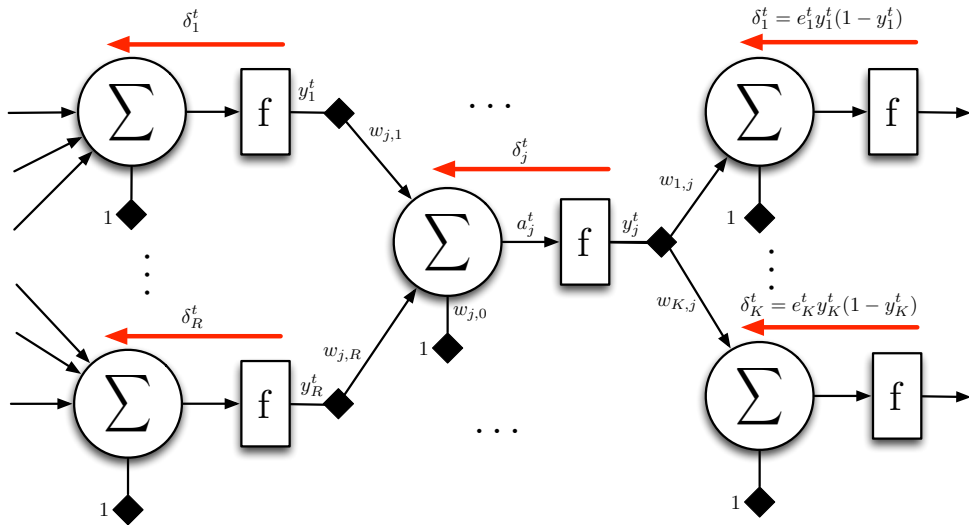
## 7.3 Error backpropagation

---

# Error backpropagation

- Learning with the multilayer perceptron: determining the weights  $\mathbf{w}, w_0$  of all neurons
- Error backpropagation
  - Learning by gradient descent
  - Output layer: error-guided correction between desired and achieved outputs
  - Hidden layers: correction according to sensitivities (influence of the neuron on the error in the output layer)

# Error backpropagation



## Neuron output values

- Value  $y_j^t$  of the neuron  $j$  for the data  $\mathbf{x}^t$

$$y_j^t = f(a_j^t) = f\left(\sum_{i=1}^R w_{j,i} y_i^t + w_{j,0}\right)$$

- $f$ : neuron activation function
- $a_j^t = \sum_{i=1}^R w_{j,i} y_i^t + w_{j,0}$ : weighted summation of neuron inputs
- $w_{j,i}$ : weight of the link connecting the neuron  $j$  to the neuron  $i$  of the previous layer
- $w_{j,0}$ : bias of the neuron  $j$
- $y_i^t$ : output of the neuron  $i$  of the previous layer for the data  $\mathbf{x}^t$
- $R$ : number of neurons on the previous layer

## Output layer error

- A dataset  $\mathcal{X} = \{\mathbf{x}^t, \mathbf{r}^t\}_{t=1}^N$ , with  $\mathbf{r}^t = [r_1^t \ r_2^t \ \dots \ r_K^t]^\top$ , where  $r_j^t = 1$  if  $\mathbf{x}^t \in C_j$ , otherwise  $r_j^t = 0$
- Error observed for data  $\mathbf{x}^t$  on neuron  $j$  of the output layer

$$e_j^t = r_j^t - y_j^t$$

- Quadratic error observed for data  $\mathbf{x}^t$  on the  $K$  neurons of the output layer (one neuron per class)

$$E^t = \frac{1}{2} \sum_{j=1}^K (e_j^t)^2$$

- Observed mean squared error for the data in dataset  $\mathcal{X}$

$$E = \frac{1}{N} \sum_{t=1}^N E^t$$

## Error correction for the output layer

- Weight correction by gradient descent of the mean squared error

$$\Delta w_{j,i} = -\eta \frac{\partial E}{\partial w_{j,i}} = -\frac{\eta}{N} \sum_{t=1}^N \frac{\partial E^t}{\partial w_{j,i}}$$

- Error of neuron  $j$  depends on the neurons of the previous layer
  - Development using the derivative chain rule ( $\frac{\partial f}{\partial x} = \frac{\partial f}{\partial y} \frac{\partial y}{\partial x}$ )

$$\begin{aligned} \frac{\partial E^t}{\partial w_{j,i}} &= \frac{\partial E^t}{\partial e_j^t} \frac{\partial e_j^t}{\partial y_j^t} \frac{\partial y_j^t}{\partial a_j^t} \frac{\partial a_j^t}{\partial w_{j,i}} \\ \frac{\partial E^t}{\partial w_{j,0}} &= \frac{\partial E^t}{\partial e_j^t} \frac{\partial e_j^t}{\partial y_j^t} \frac{\partial y_j^t}{\partial a_j^t} \frac{\partial a_j^t}{\partial w_{j,0}} \end{aligned}$$

## Calculation of partial derivatives

- Development with sigmoid activation function ( $y_j^t = \frac{1}{1+\exp(-a_j^t)}$ )

$$\frac{\partial E^t}{\partial e_j^t} = \frac{\partial}{\partial e_j^t} \frac{1}{2} \sum_{l=1}^K (e_l^t)^2 = e_j^t$$

$$\frac{\partial e_j^t}{\partial y_j^t} = \frac{\partial}{\partial y_j^t} r_j^t - y_j^t = -1$$

$$\begin{aligned} \frac{\partial y_j^t}{\partial a_j^t} &= \frac{\partial}{\partial a_j^t} \frac{1}{1 + \exp(-a_j^t)} = \frac{\exp(-a_j^t)}{[1 + \exp(-a_j^t)]^2} \\ &= \frac{1}{1 + \exp(-a_j^t)} \frac{\exp(-a_j^t) + 1 - 1}{1 + \exp(-a_j^t)} = y_j^t(1 - y_j^t) \end{aligned}$$

$$\frac{\partial a_j^t}{\partial w_{j,i}} = \frac{\partial}{\partial w_{j,i}} \sum_{l=1}^R w_{j,l} y_l^t + w_{j,0} = y_i^t$$

$$\frac{\partial a_j^t}{\partial w_{j,0}} = \frac{\partial}{\partial w_{j,0}} \sum_{l=1}^R w_{j,l} y_l^t + w_{j,0} = 1$$



## Learning for the output layer

- Learning the output layer weights

$$\begin{aligned}\Delta w_{j,i} &= -\frac{\eta}{N} \sum_{t=1}^N \frac{\partial E^t}{\partial w_{j,i}} = -\frac{\eta}{N} \sum_{t=1}^N \frac{\partial E^t}{\partial e_j^t} \frac{\partial e_j^t}{\partial y_j^t} \frac{\partial y_j^t}{\partial a_j^t} \frac{\partial a_j^t}{\partial w_{j,i}} \\ &= \frac{\eta}{N} \sum_{t=1}^N e_j^t y_j^t (1 - y_j^t) y_i^t\end{aligned}$$

- Learning the biases of the output layer

$$\begin{aligned}\Delta w_{j,0} &= -\frac{\eta}{N} \sum_{t=1}^N \frac{\partial E^t}{\partial w_{j,0}} = -\frac{\eta}{N} \sum_{t=1}^N \frac{\partial E^t}{\partial e_j^t} \frac{\partial e_j^t}{\partial y_j^t} \frac{\partial y_j^t}{\partial a_j^t} \frac{\partial a_j^t}{\partial w_{j,0}} \\ &= \frac{\eta}{N} \sum_{t=1}^N e_j^t y_j^t (1 - y_j^t)\end{aligned}$$

## 7.4 The delta rule

---

## The delta rule

- Let a delta  $\delta_j^t$ , which corresponds to the *local gradient* of the neuron  $j$  for the data  $\mathbf{x}^t$

$$\delta_j^t = e_j^t y_j^t (1 - y_j^t)$$

$$\Delta w_{j,i} = \frac{\eta}{N} \sum_{t=1}^N \delta_j^t y_i^t$$

$$\Delta w_{j,0} = \frac{\eta}{N} \sum_{t=1}^N \delta_j^t$$

- Useful formulation for hidden layer error correction

## Hidden layer error correction

- Error gradient for hidden layers

$$\frac{\partial E^t}{\partial w_{j,i}} = \frac{\partial E^t}{\partial y_j^t} \frac{\partial y_j^t}{\partial a_j^t} \frac{\partial a_j^t}{\partial w_{j,i}}$$

- Only  $\frac{\partial E^t}{\partial y_j^t}$  changes,  $\frac{\partial y_j^t}{\partial a_j^t}$  and  $\frac{\partial a_j^t}{\partial w_{j,i}}$  are the same as on the output layer
  - Error for a neuron of the hidden layer depends on the error of the  $k$  neurons of the next layer (error backpropagation)

$$E^t = \frac{1}{2} \sum_k (e_k^t)^2$$

$$\frac{\partial E^t}{\partial y_j^t} = \frac{\partial}{\partial y_j^t} \frac{1}{2} \sum_k (e_k^t)^2 = \sum_k e_k^t \frac{\partial e_k^t}{\partial y_j^t}$$

## Hidden layer error correction

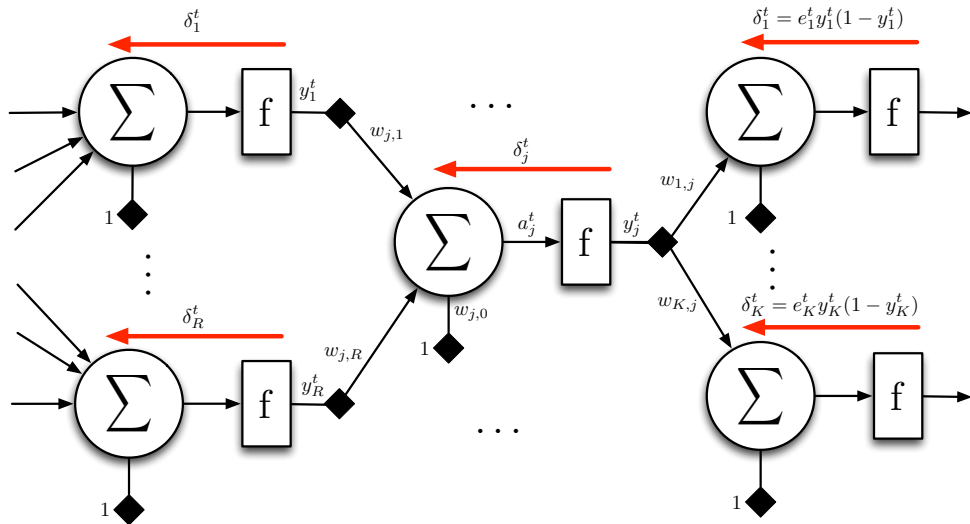
$$\begin{aligned}\frac{\partial E^t}{\partial y_j^t} &= \frac{\partial}{\partial y_j^t} \frac{1}{2} \sum_k (e_k^t)^2 = \sum_k e_k^t \frac{\partial e_k^t}{\partial y_j^t} \\&= \sum_k e_k^t \frac{\partial e_k^t}{\partial a_k^t} \frac{\partial a_k^t}{\partial y_j^t} \\&= \sum_k e_k^t \frac{\partial (r_k^t - y_k^t)}{\partial a_k^t} \frac{\partial (\sum_l w_{k,l} y_l^t + w_{k,0})}{\partial y_j^t} \\&= \sum_k e_k^t [-y_k^t (1 - y_k^t)] w_{k,j} \\ \delta_k^t &= e_k^t [y_k^t (1 - y_k^t)] \\ \frac{\partial E^t}{\partial y_j^t} &= - \sum_k \delta_k^t w_{k,j}\end{aligned}$$

## Hidden layer error correction

- Correction of the corresponding error

$$\begin{aligned}\frac{\partial E^t}{\partial w_{j,i}} &= \frac{\partial E^t}{\partial y_j^t} \frac{\partial y_j^t}{\partial a_j^t} \frac{\partial a_j^t}{\partial w_{j,i}} \\ &= - \left[ \sum_k \delta_k^t w_{k,j} \right] y_j^t (1 - y_j^t) y_i^t \\ \delta_j^t &= y_j^t (1 - y_j^t) \sum_k \delta_k^t w_{k,j} \\ \Delta w_{j,i} &= -\eta \frac{\partial E}{\partial w_{j,i}} = -\frac{\eta}{N} \sum_{t=1}^N \frac{\partial E^t}{\partial w_{j,i}} = \frac{\eta}{N} \sum_{t=1}^N \delta_j^t y_i^t \\ \Delta w_{j,0} &= -\eta \frac{\partial E}{\partial w_{j,0}} = -\frac{\eta}{N} \sum_{t=1}^N \frac{\partial E^t}{\partial w_{j,0}} = \frac{\eta}{N} \sum_{t=1}^N \delta_j^t\end{aligned}$$

# Error backpropagation



## 7.5 Backpropagation algorithm

---



# Batch and online learning

- Batch learning
  - Guided by mean squared error ( $E = \frac{1}{N} \sum_t E^t$ )
  - Weight correction once at each epoch, calculating the error for the whole dataset
  - Relatively stable learning
- Online learning
  - Weight correction for each data presentation, so  $N$  weight corrections per epoch
  - Guided by the quadratic error of each data ( $E^t$ )
  - Requires permutation of the processing order at each epoch to avoid bad sequences
  - Online learning faster than batch, but with greater instabilities
- Mini-batch learning
  - Trade-off between online learning and batch learning, using mini batches of a predefined size

# Neuron saturation

- Operating range of neurons with sigmoid function around 0
  - For low  $a$  values  $f_{sig}(a) \rightarrow 0$ , and for high  $a$  values,  $f_{sig}(a) \rightarrow 1$

$$f_{sig}(1) = 0.7311, \quad f_{sig}(5) = 0.9933, \quad f_{sig}(10) \approx 1$$

- For large/small values, say  $x < -10$  or  $x > 10$ , gradient almost equal to zero
  - Extremely slow learning
- Input values, the  $\mathbf{x}^t$ , must be normalized beforehand in  $[-1, 1]$ 
  - Typically, normalization according to min and max values of the dataset for each dimension
  - Apply the same normalization to the evaluated data (do not recalculate the normalization)

## Target output values

- In classification, target values  $r_i^t \in \{0, 1\}$ 
  - Also suffers from the problem of neuron saturation with sigmoid function
  - We aim to approximate the  $r_i^t$  with the neurons of the output layer

$$f_{sig}(a) = 0 \Rightarrow a \rightarrow -\infty, \quad f_{sig}(a) = 1 \Rightarrow a \rightarrow \infty$$

- Solution: transform the desired values into values  $\tilde{r}_i^t \in \{0.05, 0.95\}$ 
  - If  $\mathbf{x}^t \in C_i$  then  $\tilde{r}_i^t = 0.95$
  - Otherwise  $\tilde{r}_i^t = 0.05$

# Weights initialization

- The weights and biases of a multilayer perceptron are randomly initialized
  - Typically, weights and biases are initialized uniformly in  $[-0.5, 0.5]$

$$w_{j,i} \sim \mathcal{U}(-0.5, 0.5), \forall i, j$$

- Multilayer Perceptron is thus a stochastic algorithm
  - From one run to another, we do not necessarily obtain the same results

# Backpropagation algorithm

1. Normalize data  $x_i^t \in [-1,1]$  and target output  $\tilde{r}_j^t \in \{0.05, 0.95\}$
2. Initialize weights and bias randomly,  $w_{i,j} \in [-0.5, 0.5]$
3. As long as the stop criterion is not reached, repeat:
  - 3.1 Calculate the observed outputs by propagating the data forward
  - 3.2 Calculate the observed errors on the output layer

$$e_j^t = \tilde{r}_j^t - y_j^t, \quad j = 1, \dots, K, \quad t = 1, \dots, N$$

- 3.3 Adjust weights and bias by backpropagating the observed error

$$w_{j,i} = w_{j,i} + \Delta w_{j,i} = w_{j,i} + \frac{\eta}{N} \sum_t \delta_j^t y_i^t$$

$$w_{j,0} = w_{j,0} + \Delta w_{j,0} = w_{j,0} + \frac{\eta}{N} \sum_t \delta_j^t$$

where the local gradient is defined by:


$$\delta_j^t = \begin{cases} e_j^t y_j^t (1 - y_j^t) & \text{if } j \in \text{output layer} \\ y_j^t (1 - y_j^t) \sum_k \delta_k^t w_{k,j} & \text{if } j \in \text{hidden layer} \end{cases}$$

## **7.6 Training techniques and tips**

---

# Overfitting and stop criterion

- Number of epochs: determining factor for overfitting
- Stop criterion: when the error on the validation set increases (generalization)
- Requires to use part of the dataset for validation



# Momentum

- Generalized delta rule


$$w_{j,i}(n) = w_{j,i}(n-1) + \frac{\eta}{N} \sum_t \delta_j^t y_i^t + \alpha \Delta w_{j,i}(n-1)$$

$$w_{j,0}(n) = w_{j,0}(n-1) + \frac{\eta}{N} \sum_t \delta_j^t + \alpha \Delta w_{j,0}(n-1)$$

- Factor  $\Delta w_{j,i}(n-1)$  is the correction made to the weight/bias at the previous epoch
- Parameter  $\alpha \in [0.5, 1]$  is named *momentum*
- Gives inertia to the descent of the gradient, including a correction from the previous iterations
- With momentum, the factor  $\Delta w_{j,i}(n-1)$  depends itself on the correction of the previous iteration  $\Delta w_{j,i}(n-2)$ , and so on



# Momentum



# Regression with multilayer perceptron

- Backpropagation algorithm developed for sigmoid transfer function for classification
  - Other transfer functions can be used
    - Linear function:  $f_{lin}(a) = a$
    - Hyperbolic tangent function:  $f_{tanh}(a) = \tanh(a)$
    - ReLU function (*rectified linear unit*):  $f_{ReLU}(a) = \max(0, a)$
  - In fact, all continuous functions derivable on  $\mathbb{R}$  can be used
- Multilayer perceptron suitable for regression
  - Recommended topology: a hidden layer with a sigmoid function and an output layer with a linear function
  - Mean squared error criterion is appropriate for regression

## Second order method

- The gradient descent is a first order method (first derivatives)
- Possibility to do better with second order methods
- Newton's method
  - Based on the expansion of the second order Taylor series,  $\mathbf{x}' = \mathbf{x} + \Delta\mathbf{x}$  one point in the neighbourhood of  $\mathbf{x}$

$$F(\mathbf{x}') = F(\mathbf{x} + \Delta\mathbf{x}) \approx F(\mathbf{x}) + \nabla F(\mathbf{x})^\top \Delta\mathbf{x} + \frac{1}{2} \Delta\mathbf{x}^\top \nabla^2 F(\mathbf{x}) \Delta\mathbf{x} = \hat{F}(\mathbf{x})$$

- Search for a plateau in the squared error  $\hat{F}(\mathbf{x})$

$$\begin{aligned} \frac{\partial \hat{F}(\mathbf{x})}{\partial \mathbf{x}} &= \nabla F(\mathbf{x}) + \nabla^2 F(\mathbf{x}) \Delta\mathbf{x} = 0 \\ \Delta\mathbf{x} &= -(\nabla^2 F(\mathbf{x}))^{-1} \nabla F(\mathbf{x}) \end{aligned}$$

- Calculation of the inverse of the Hessian matrix  $((\nabla^2 F(\mathbf{x}))^{-1})$ : high calculation costs
- Conjugate gradient method avoids the calculation of the inverse of the Hessian matrix

## 7.7 Multilayer perceptron in scikit-learn

---

- Multilayer perceptron is available in scikit-learn
  - Scikit-learn uses some (but not all) of the deep network advances
  - No GPU acceleration for calculations, rigid models (not easily customizable)
- `neural_network.MLPClassifier`: multilayer perceptron for classification
  - Minimizes cross entropy for classification with gradient-based methods

$$E_{entr} = - \sum_t r^t \log y^t + (1 - r^t) \log(1 - y^t)$$

- `neural_network.MLPRegressor`: multilayer perceptron for regression
  - Minimizes the quadratic error with gradient-based methods

## MLPClassifier and MLPRegressor parameters

- `hidden_layer_sizes` (tuple): number of neurons on each hidden layer (default: (100,))
- `activation` (string): 'identity' (linear), 'logistic' (sigmoid), 'tanh' and 'relu' (default: 'relu')
- `solver` (string): 'lbfgs' (quasi-Newton), 'sgd' (stochastic gradient descent), 'adam' (sgd with automatic determination of the learning rate) (default: 'adam')
- `alpha` (float): parameter of the  $L_2$  regulation of the weights (default: 0.0001)
- `batch_size` (int): batch size for each update (default:  $\min(200, N)$ )
- `learning_rate_init` (float): initial learning rate (default: 0.001)
- `learning_rate` (string): 'constant', 'invscaling' ( $\text{learning\_rate\_init} / \text{pow}(t, \text{power\_t})$ ), 'adaptive' (reduces current rate when learning stagnates) (default: 'constant')
- `max_iter` (int): maximum number of epochs (default: 200)
- `tol` (float): tolerance, stop learning if gain  $<$  tolerance for more than two epochs (default:  $10^{-4}$ )
- `momentum` (float): momentum for gradient descent (default: 0.9)
- `early_stopping` (bool): stop when error on validation set does not go down anymore (default: False)
- `validation_fraction` (float): portion of the data used for validation with the *early stopping*. (default: 0.1)